

Lecture 6: Linear Regression with One Regressor

Zheng Tian

Contents

1	Introduction	1
2	The Linear Regression Model	2
3	The OLS Estimation Method for a Linear Regression Model	5
4	Algebraic Properties of the OLS Estimator	12
5	Measures of Fit	14
6	The Least Squares Assumptions	16
7	Sampling Distribution of the OLS Estimators	18

1 Introduction

This lecture introduces a linear regression model with one regressor called a simple linear regression model. We will learn the ordinary least squares (OLS) method to estimate a simple linear regression model, discuss the algebraic and statistical properties of the OLS estimator, introduce two measures of goodness of fit, and bring up three least squares assumptions for a linear regression model. As an example, we apply the OLS estimation method to a linear model of test scores and class sizes in California school districts.

This lecture lays out foundations for all lectures to come. Although in practice we seldom use a linear regression model with only one regressor, the essential principles of the OLS estimation method and hypothesis testing are the same for a linear regression model with multiple regressors.

2 The Linear Regression Model

2.1 What is regression?

Definition of regress in Merriam-Webster's dictionary

Merriam-Webster gives the following definition of the word "regress":

1. An act or the privilege of going or coming back
2. Movement backward to a previous and especially worse or more primitive state or condition
3. The act of reasoning backward

The meaning of regression in statistics?

In statistical modeling, regression analysis is a statistical process for estimating the relationships among variables.¹ Specifically, most regression analysis focus on the conditional mean of the dependent variable given the independent variables, which is a function of the values of independent variables.

A very simple functional form of a conditional expectation is a linear function. That is, we can model the conditional mean as follows,

$$E(Y | X = x) = f(x) = \beta_0 + \beta_1 x \quad (1)$$

Equation 1 is called a **simple linear regression function**.

2.2 An example: Test scores versus class size

Let's go back to the example of California school districts, introduced in Lecture 1.

Research question:

The research question of this application is: Can reducing class size increase students' test scores?

How can we answer the question?

- We randomly choose 42 students and divide them into two classes, with one having 20 students and another having 22. And they are taught with the same subject and by the same teachers.

¹Wikipedia, the free encyclopedia. Regression analysis. Retrieved from https://en.wikipedia.org/wiki/Regression_analysis

- Randomization ensures that it is the difference in class sizes of the two classes that is the only factor affecting test scores.
- After a test for both classes, we then compute the average test scores that can be expressed as,

$$\begin{aligned} E(\text{TestScore}|\text{ClassSize} = 20) \\ E(\text{TestScore}|\text{ClassSize} = 22) \end{aligned}$$

- Then the effect of class size on test scores is the difference in the conditional means, i.e.,

$$E(\text{TestScore}|\text{ClassSize} = 20) - E(\text{TestScore}|\text{ClassSize} = 22)$$

- If the difference is large enough, we can say that reducing class can improve students' test performance.

A simple linear regression model of test scores v.s. class size

As mentioned above, a simple linear regression function can be used to describe the relationship between test scores and class sizes. Since it regards the association between these two variable for the whole population, we call this regression function as the **population regression function** or the **population regression line**, taking the following form,

$$E(\text{TestScore}|\text{ClassSize}) = \beta_0 + \beta_1 \text{ClassSize} \quad (2)$$

By calculating the conditional expectation, some other factors, apart from class sizes, are left out of the population regression function. Although these factors may also influence test scores, they are either unimportant in your reasoning or unable to be measured. We can lump all these factors into a single term, and set up a **simple linear regression model** as follows,

$$\text{TestScore} = \beta_0 + \beta_1 \text{ClassSize} + \text{OtherFactors} \quad (3)$$

If we assume $E(\text{OtherFactors}|\text{ClassSize}) = 0$, then the simple linear regression model (Eq. 3) becomes the population regression line (Eq. 2).

A distinction between the population regression function and the population regression model

Note that here we have two concepts: the population regression function and the population regression model. What's their difference? Simply put,

- A population regression function gives us a deterministic relation between class size and the expectation of test scores. That is, when we have a value of class size and know the values of β_0 and β_1 , there is one and only one expected value of test scores associated with this class size. However, we cannot compute the exact value of the test score of a particular observation.
- A population regression model, by including other factors, gives us a complete description of a data generating process (DGP). That is, when we have all the values of class sizes and other factors and know β_0 and β_1 , we can generate all the values of test scores. Also, when we consider other factors as a random variable, the association between test scores and class size is not deterministic, depending on the value of other factors.

An interpretation of the population regression model

Now we have set up the simple linear regression model,

$$TestScore = \beta_0 + \beta_1 ClassSize + OtherFactors$$

What is β_1 and β_0 represent in the model?

- Interpret β_1

Let's first look at β_1 . When we hold other factors constant, the only reason for a change in test scores is a change in class size. Denote $\Delta TestScore$ and $\Delta ClassSize$ to be their respective change. According to the above regression model, holding other factors constant, we have

$$\Delta TestScore = \beta_1 \Delta ClassSize$$

where β_0 is removed because it is also a constant. Then, we get

$$\beta_1 = \frac{\Delta TestScore}{\Delta ClassSize}$$

That is, β_1 measures the change in the test score resulting from a **one-unit change** in the class size. When $TestScore$ and $ClassSize$ are two continuous variable, we can write β_1 as

$$\beta_1 = \frac{dTestScore}{dClassSize}$$

Hence, we often call β_1 as the **marginal effect** of the class size on the test score.

The phrase of "holding other factors constant" is important. Without it, we cannot disentangle the effect of class sizes on test scores from other factors. "Holding other things constant" is often expressed as the notion of **ceteris paribus**.

- Interpret β_0

β_0 is the intercept in the model. Sometimes it bears real meanings, but sometimes it merely presents as an intercept. In this regression model, β_0 is the test score when the

class size and other factors are all zero, which is obviously nonsensical. Thus, β_0 does not have a real meaning in this model, and it just determines where the population regression line intersects the Y axis.

2.3 The general linear regression model

Let's generalize test scores and class sizes to be two random variables Y and X . For both, there are n observations so that each observation $i = 1, 2, 3, \dots$ is associated with a pair of values of (X_i, Y_i) .

Then a **simple linear regression model** that associates Y with X is

$$Y_i = \beta_0 + \beta_1 X_i + u_i, \text{ for } i = 1, \dots, n \quad (4)$$

- Y_i is called the dependent variable, the regressand, or the LHS (left-hand side) variable.
- X_i is called the independent variable, the regressor, or the RHS (right-hand side) variable.
- β_0 is the intercept, or the constant term. It can either have economic meaning or have merely mathematical sense, which determines the level of the regression line, i.e., the point of intersection with the Y axis.
- β_1 is the slope of the population regression line. Since $\beta_1 = dY_i/dX_i$, it is the marginal effect of X on Y . That is, holding other things constant, one unit change in X will make Y change by β_1 units.
- u_i is the error term. $u_i = Y_i - (\beta_0 + \beta_1 X_i)$ incorporates all the other factors besides X that determine the value of Y .
- $\beta_0 + \beta_1 X_i$ represents the population regression function (or the population regression line).

2.4 An graphical illustration of a linear regression model

The relationship between the data points, the population regression line, and the errors (other factors) are illustrated in Figure 1.

3 The OLS Estimation Method for a Linear Regression Model

3.1 The intuition for the OLS and minimization

The most commonly used method to estimate a linear regression model, like Equation 4, is the ordinary least squares (OLS) estimation.

Let's explain the basic idea of the OLS by dissecting its name.




Figure 1: The Population Regression Line

Ordinary It means that the OLS estimator is a very basic method, from which we may derive some variations of the OLS estimator, such as the weighted least squares (WLS), and the generalized least squares (GLS).

Least It means that the OLS estimator tries to minimize something. The "something" is the mistakes we make when we try to guess (estimate) the values of the parameters in the model. From Equation 4, if our guess for β_0 and β_1 is b_0 and b_1 , then the mistake of our guess is $\hat{u}_i = Y_i - b_0 - b_1 X_i$.

Squares It represent the actual thing (a quantity) that we minimize. The OLS does not attempt to minimize each \hat{u}_i but to minimize the sum of the squared mistakes, $\sum_{i=1}^n \hat{u}_i^2$. Taking square is to avoid possible offsetting between positive and negative values of \hat{u}_i in $\sum_i \hat{u}_i$.

3.2 The OLS estimators for β_0 and β_1

Let b_0 and b_1 be some estimators of β_0 and β_1 , respectively.² Then, the OLS estimator is the solution to the following minimization problem.

$$\min_{b_0, b_1} S(b_0, b_1) = \sum_{i=1}^n \hat{u}_i^2 = \sum_{i=1}^n (Y_i - b_0 - b_1 X_i)^2 \quad (5)$$

²Recall that an **estimator** is a function of a sample of data. An **estimate** is the numerical value of the estimator when it is computed using data from a sample.

where $S(b_0, b_1)$ is a function of b_0 and b_1 , measuring the sum of the squared prediction mistakes over all n observation.

The mathematical derivation of the OLS estimators for β_0 and β_1

We solve the problem by taking the derivative of $S(b_0, b_1)$ with respect to b_0 and b_1 , respectively. Suppose $b_0^* = \hat{\beta}_0$ and $b_1^* = \hat{\beta}_1$ are the solution to the minimization problem. Then the first order conditions evaluated at $(\hat{\beta}_0, \hat{\beta}_1)$ are

$$\frac{\partial S}{\partial b_0}(\hat{\beta}_0, \hat{\beta}_1) = \sum_{i=1}^n (-2)(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = 0 \quad (6)$$

$$\frac{\partial S}{\partial b_1}(\hat{\beta}_0, \hat{\beta}_1) = \sum_{i=1}^n (-2)(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)X_i = 0 \quad (7)$$

Rearranging Equation 6, we get

$$\begin{aligned} \sum_{i=1}^n Y_i - n\hat{\beta}_0 - \hat{\beta}_1 \sum_{i=1}^n X_i &= 0 \\ \hat{\beta}_0 &= \frac{1}{n} \sum_{i=1}^n Y_i - \frac{\hat{\beta}_1}{n} \sum_{i=1}^n X_i = \bar{Y} - \hat{\beta}_1 \bar{X} \end{aligned} \quad (8)$$

Rearranging Equation 7 and plugging Equation 8, we get

$$\begin{aligned} \sum_{i=1}^n X_i Y_i - \hat{\beta}_0 \sum_{i=1}^n X_i - \hat{\beta}_1 \sum_{i=1}^n X_i^2 &= 0 \\ \sum_{i=1}^n X_i Y_i - \frac{1}{n} \sum_{i=1}^n X_i \sum_{i=1}^n Y_i + \hat{\beta}_1 \frac{1}{n} \left(\sum_{i=1}^n X_i \right)^2 - \hat{\beta}_1 \sum_{i=1}^n X_i^2 &= 0 \\ \hat{\beta}_1 &= \frac{n \sum_{i=1}^n X_i Y_i - \sum_{i=1}^n X_i \sum_{i=1}^n Y_i}{n \sum_{i=1}^n X_i^2 - (\sum_{i=1}^n X_i)^2} \end{aligned} \quad (9)$$

For the numerator in Equation 9, we can show the following

$$\begin{aligned} \sum_i (X_i - \bar{X})(Y_i - \bar{Y}) &= \sum_i X_i Y_i - \bar{X} \sum_i Y_i - \bar{Y} \sum_i X_i + \sum_i \bar{X} \bar{Y} \\ &= \sum_i X_i Y_i - 2n\bar{X}\bar{Y} + n\bar{X}\bar{Y} \\ &= \sum_i X_i Y_i - n\bar{X}\bar{Y} \\ &= \frac{1}{n} \left(n \sum_i X_i Y_i - \sum_i X_i \sum_i Y_i \right) \end{aligned}$$

Similarly, we can show that $\sum_i (X_i - \bar{X})^2 = \frac{1}{n} [n \sum_i X_i^2 - (\sum_i X_i)^2]$.

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

Since we know that the sample covariance of X and Y is $s_{XY} = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})$ and the sample variance of X is $s_X^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$, the equation above can also be written as

$$\hat{\beta}_1 = \frac{s_{XY}}{s_X^2}$$

In sum, solving the minimization problem (Equation 5), we obtain the OLS estimators for β_0 and β_1 as

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{s_{XY}}{s_X^2} \quad (10)$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X} \quad (11)$$

3.3 The predicted values, residuals, and the sample regression line

The predicted values

- After obtaining the estimators, we can compute the **predicted values** \hat{Y}_i for $i = 1, \dots, n$

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$$

- The line represented by the above equation is called **the sample regression line**.
- The sample average point (\bar{X}, \bar{Y}) is always on the sample regression line because, from Equation 11, we have

$$\bar{Y} = \hat{\beta}_0 + \hat{\beta}_1 \bar{X}$$

The residuals

- The **residuals** \hat{u}_i for $i = 1, \dots, n$ are

$$\hat{u}_i = Y_i - \hat{Y}_i$$

- The residuals are the difference between the observed values of Y_i and its predicted value. That is, they are the actual prediction errors we make when using the OLS estimators.

3.4 A comparison between the population regression model and the sample counterparts

We should pause here to make a clear distinction between the population regression function and model and their counterparts.

The population regression function versus the sample regression function

- The population regression function is a function between the conditional mean of Y given X and X , that is,

$$E(Y | X) = \beta_0 + \beta_1 X_i$$

where β_0 and β_1 are the population parameters.

- The sample regression function is a function between the predicted value and X , that is,

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$$

The regression errors versus residuals

- The error term, u_i , in the population regression model represents the other factors that the population regression function does not take into account. It is the difference between Y_i and $E(Y_i | X_i)$. Thus, we have

$$Y_i = \beta_0 + \beta_1 X_i + u_i$$

- The residuals, \hat{u}_i , represent the actual mistakes we make with a set of estimators. It is the difference between Y_i and its predicted value \hat{Y}_i . Thus, we have

$$Y_i = \hat{\beta}_0 + \hat{\beta}_1 X_i + \hat{u}_i$$

Table 1: A comparison between the population regression and its sample counterparts

	Population	Sample
Regression functions	$\beta_0 + \beta_1 X_i$	$\hat{\beta}_0 + \hat{\beta}_1 X_i$
Parameters	β_0, β_1	$\hat{\beta}_0, \hat{\beta}_1$
Errors vs residuals	u_i	\hat{u}_i
The regression model	$Y_i = \beta_0 + \beta_1 X_i + u_i$	$Y_i = \hat{\beta}_0 + \hat{\beta}_1 X_i + \hat{u}_i$

3.5 The OLS estimates of the relationship between test scores and the student-teacher ratio

Let's come back to the application of test scores versus the student-teacher ratios in California school districts. The goal is to estimate the effect of class sizes, measured by the student-teacher

ratios, on test scores. Before setting up a formal regression model, it is always a good practice to glance over the data using some exploratory data analysis techniques.

Exploratory analysis

- Basic summary statistics

We first need to compute basic summary statistics to see the sample distribution of the data. Some commonly used summary statistics include mean, standard deviation, median, minimum, maximum, and quantile (percentile), etc. Table 2 summarizes the distribution of test scores and class sizes for the sample.

Table 2: Summary Of distributions of student-teacher ratios and test scores

	Average	S.t.d.	10%	25%	40%	50%	60%	75%	90%
<i>TestScore</i>	654.16	19.05	630.4	640.05	649.07	654.45	659.4	666.66	678.86
<i>STR</i>	19.64	1.89	17.35	18.58	19.27	19.72	20.08	20.87	21.87

- Scatterplot

A scatterplot visualizes the relationship between two variables straightforwardly, which is helpful for us to decide what a functional form a regression model should properly take. Figure 2 shows that test scores and student-teacher ratios may be negatively related. The correlation coefficient between the two variables is -0.23, verifying the existence of a weak negative relationship.




Figure 2: The scatterplot between student-teacher ratios and test scores

Regression analysis

After exploratory analysis, we can estimate the linear model. Although the formula of computing β_1 and β_0 (Equations 10 and 11) seems complicated, the practical estimation procedure is simplified by using computer software, like R. For now, let's simply present the estimation results in the following equation,

$$\widehat{TestScore} = 698.93 - 2.28 \times STR \quad (12)$$

We can draw the sample regression line represented by Equation 12 in the scatterplot to eyeball how well the regression model fits the data.




Figure 3: The estimated regression line for the California data

Interpretation of the estimated coefficients

Upon obtaining the coefficient estimates, what we need to do next includes hypothesis tests, model specification tests, robustness (or sensitivity) test, and interpretation. Let's first see how to correctly interpret the estimation results.

- Our main interest is in the slope that tell us how much a unit change in student-teacher ratios will cause test scores to change. The slope of -2.28 means that an increase in the student-teacher ratio by one student per class is, on average, associated with a decline in district-wide test scores by 2.28 points on the test.
- The intercept literally means that if the student-teacher ratio is zero, the average district-wide test scores will be 698.9. However, it is nonsense for having some positive test scores when the student-teacher ratio is zero. Therefore, the intercept term in this case merely serves as determining the level of the sample regression line.

- The mere number of -2.28 really does not make much sense for the readers of your research. We have to put it into the context of California school district to avoid ridiculous results even though the estimation itself is correct. (Read the discussion in the paragraphs in Page 117.)

4 Algebraic Properties of the OLS Estimator

The OLS estimator has many good properties. Let's first look at some of its algebraic properties. That is, these properties are the results of the minimization problem in Equation (5), regardless of any statistical assumptions we will introduce in the next sections.

4.1 TSS, ESS, and SSR

- From $Y_i = \hat{Y}_i + \hat{u}_i$, we can define
 - **The total sum of squares:** $TSS = \sum_{i=1}^n (Y_i - \bar{Y})^2$
 - **The explained sum of squares:** $ESS = \sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2$
 - **The sum of squared residuals:** $SSR = \sum_{i=1}^n (Y_i - \hat{Y}_i)^2 = \sum_{i=1}^n \hat{u}_i^2$

Note that TSS, ESS, and SSR all take the form of "deviation from the mean". This form is only valid when an intercept is included in the regression model.³

4.2 Some algebraic properties among \hat{u}_i , \hat{Y}_i , Y_i , and X_i

The OLS residuals and the predicted values satisfy the following equations:⁴

$$\sum_{i=1}^n \hat{u}_i = 0 \quad (13)$$

$$\frac{1}{n} \sum_{i=1}^n \hat{Y}_i = \bar{Y} \quad (14)$$

$$\sum_{i=1}^n \hat{u}_i X_i = 0 \quad (15)$$

$$TSS = ESS + SSR \quad (16)$$

³We are not going to prove this because it involves higher level knowledge of linear algebra. You can estimate a linear regression model of $Y_i = \beta_1 X_i + u_i$, for which TSS is simply $\sum_i Y_i^2$ and ESS is $\sum_i \hat{Y}_i^2$. Also, for this model, $\sum_i \hat{u}_i \neq 0$.

⁴Equation 13 holds only for a linear regression model with an intercept, but Equation 15 holds regardless of the presence of an intercept.

4.3 The proof of these properties

Here, I just prove Equation 16. The proofs for the other equations above are in Appendix 4.3 in the textbook.

Proof of Equation 13

From Equation 11 we can write the OLS residuals as

$$\hat{u}_i = Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i = (Y_i - \bar{Y}) - \hat{\beta}_1 (X_i - \bar{X})$$

Thus

$$\sum_{i=1}^n \hat{u}_i = \sum_{i=1}^n (Y_i - \bar{Y}) - \hat{\beta}_1 \sum_{i=1}^n (X_i - \bar{X})$$

By definition of the sample average, we have

$$\sum_{i=1}^n (Y_i - \bar{Y}) = 0 \text{ and } \sum_{i=1}^n (X_i - \bar{X}) = 0$$

It follows that $\sum_{i=1}^n \hat{u}_i = 0$.

Proof of Equation 14

Note that $Y_i = \hat{Y}_i + \hat{u}_i$. So

$$\sum_{i=1}^n Y_i = \sum_{i=1}^n \hat{Y}_i + \sum_{i=1}^n \hat{u}_i = \sum_{i=1}^n \hat{Y}_i$$

It follows that $\bar{\hat{Y}} = (1/n) \sum_{i=1}^n \hat{Y}_i = \bar{Y}$.

Proof of Equation 15

$\sum_{i=1}^n \hat{u}_i = 0$ implies that

$$\begin{aligned} & \sum_{i=1}^n \hat{u}_i X_i \\ &= \sum_{i=1}^n \hat{u}_i (X_i - \bar{X}) \\ &= \sum_{i=1}^n \left[(Y_i - \bar{Y}) - \hat{\beta}_1 (X_i - \bar{X}) \right] (X_i - \bar{X}) \\ &= \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y}) - \hat{\beta}_1 \sum_{i=1}^n (X_i - \bar{X})^2 = 0 \end{aligned}$$

Proof of $TSS = ESS + SSR$

$$\begin{aligned}
TSS &= \sum_{i=1}^n (Y_i - \bar{Y})^2 = \sum_{i=1}^n (Y_i - \hat{Y}_i + \hat{Y}_i - \bar{Y})^2 \\
&= \sum_{i=1}^n (Y_i - \hat{Y}_i)^2 + \sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2 + 2 \sum_{i=1}^n (Y_i - \hat{Y}_i)(\hat{Y}_i - \bar{Y}) \\
&= SSR + ESS + 2 \sum_{i=1}^n \hat{u}_i \hat{Y}_i \\
&= SSR + ESS + 2 \sum_{i=1}^n \hat{u}_i (\hat{\beta}_0 + \hat{\beta}_1 X_i) \\
&= SSR + ESS + 2(\hat{\beta}_0 \sum_{i=1}^n \hat{u}_i + \hat{\beta}_1 \sum_{i=1}^n \hat{u}_i X_i) \\
&= SSR + ESS
\end{aligned}$$

where the final equality follows from Equations 13 and 15.

5 Measures of Fit

5.1 Goodness of Fit: R^2

R^2 is one of the commonly used measures for how well the OLS regression line fits the data. R^2 is the fraction of the sample variance of Y_i explained by X_i . The sample variance can be represented with TSS and the part of sample variance explained by X can be represented by ESS . Therefore, mathematically, we can define R^2 as

$$R^2 = \frac{ESS}{TSS} = 1 - \frac{SSR}{TSS} \quad (17)$$

R^2 is often called the coefficient of determination. It indicates the proportion of the variance in the dependent variable that is predictable from the independent variable(s).

Properties of R^2

- $R^2 \in [0, 1]$

$R^2 = 0$ when $\hat{\beta}_1 = 0$, that is, X cannot explain the variance in Y .

$$\hat{\beta}_1 = 0 \Rightarrow Y_i = \hat{\beta}_0 + \hat{u}_i \Rightarrow \hat{Y}_i = \bar{Y} = \hat{\beta}_0 \Rightarrow ESS = \sum_i^n (\hat{Y}_i - \bar{Y})^2 = 0 \Rightarrow R^2 = 0$$

$R^2 = 1$ when $\hat{u}_i = 0$ for all $i = 1, \dots, n$, that is, the regression line fits all the sample data

perfectly.

$$\hat{u}_i = 0 \Rightarrow SSR = \sum_i^n \hat{u}_i^2 = 0 \Rightarrow R^2 = 1$$

- $R^2 = r_{XY}^2$

r_{XY} is the sample correlation coefficient, that is,

$$r_{XY} = \frac{S_{XY}}{S_X S_Y} = \frac{\sum_i^n (X_i - \bar{X})(Y_i - \bar{Y})}{[\sum_i^n (X_i - \bar{X})^2 \sum_i^n (Y_i - \bar{Y})^2]^{1/2}}$$

To prove $R^2 = r_{XY}^2$, let's look at SSR .

$$\begin{aligned} SSR &= \sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2 = \sum_{i=1}^n (\hat{\beta}_0 + \hat{\beta}_1 X_i - \bar{Y})^2 \\ &= \sum_{i=1}^n (\bar{Y} - \hat{\beta}_1 \bar{X} + \hat{\beta}_1 X_i - \bar{Y})^2 \\ &= \sum_{i=1}^n [\hat{\beta}_1 (X_i - \bar{X})]^2 = \hat{\beta}_1^2 \sum_{i=1}^n (X_i - \bar{X})^2 \\ &= \left[\frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]^2 \sum_{i=1}^n (X_i - \bar{X})^2 \\ &= \frac{[\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})]^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \end{aligned}$$

It follows that

$$R^2 = \frac{SSR}{TSS} = \frac{[\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})]^2}{\sum_{i=1}^n (X_i - \bar{X})^2 \sum_{i=1}^n (Y_i - \bar{Y})^2} = r_{XY}^2$$

Note: This property holds only for the linear regression model with **one regressor and an intercept**.

The use of R^2

- R^2 is usually the first statistics that we look at for judging how well the regression model fits the data.
- Most computer programs for econometrics and statistics report R^2 in their estimation results.
- However, we cannot merely rely on R^2 for judge whether the regression model is "good" or "bad". For that, we have to use some statistics that will be taught soon.

5.2 The standard error of regression (SER) as a measure of fit

Like R^2 , the standard error of regression (SER) is another measure of fit for the OLS regression.

$$\text{SER} = \sqrt{\frac{1}{n-2} \sum_{i=1}^n \hat{u}_i^2} = s \quad (18)$$

- SER has the same unit of u_i , which are the unit of Y_i .
- SER measures the average “size” of the OLS residual (the average “mistake” made by the OLS regression line).
- The root mean squared error (RMSE) is closely related to the SER:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=2}^n \hat{u}_i^2}$$

As $n \rightarrow \infty$, $\text{SER} = \text{RMSE}$.

5.3 R^2 and SER for the application of test scores v.s. class sizes

- In the application of test scores v.s. class sizes, R^2 is 0.051 or 5.1%, which implies that the regressor STR explains only 5.1% of the variance of the dependent variable $TestScore$.
- SER is 18.6, which means that standard deviation of the regression residuals is 18.6 points on the test. The magnitude of SER is so large that, in another way, shows that the simple linear regression model does not fit the data well.

6 The Least Squares Assumptions

The last two sections regard the algebraic properties of the OLS estimators. Now let's turn to their statistical properties, which are built on the following assumptions.

6.1 Assumption 1: The conditional mean of u_i given X_i is zero

$$E(u_i|X_i) = 0 \quad (19)$$

If Equation 19 is satisfied, then X_i is called **exogenous**. This assumption can be stated a little stronger as $E(u|X = x) = 0$ for any value x , that is $E(u_i|X_1, \dots, X_n) = 0$.

Since $E(u|X = x) = 0$, it follows that $E(u) = E(E(u|X)) = E(0) = 0$. The unconditional mean of u is also zero.

- A benchmark for thinking about this assumption is to consider an ideal randomized controlled experiment.

Because X is assigned randomly, all other individual characteristics – the things that make up u – are distributed independently of X , so u and X are independent. Thus, in an ideal randomized controlled experiment, $E(u|X = x) = 0$.

- In actual experiments, or with observational data, we will need to think hard about whether $E(u|X = x) = 0$ holds.

Assumption 1 can be illustrated by Figure 4. The conditional mean, $E(Y | X)$, of the conditional density distribution, $f(y | x)$, is vertically projected right on the population regression line $\beta_0 + \beta_1 X$ because $E(Y | X) = \beta_0 + \beta_1 X + E(u | X) = \beta_0 + \beta_1 X$.




Figure 4: An illustration of $E(u|X = x) = 0$

- Correlation and conditional mean

$$E(u_i|X_i) = 0 \Rightarrow \text{Cov}(u_i, X_i) = 0$$

That is, the zero conditional mean of u_i given X_i means that they are uncorrelated.

$$\begin{aligned} \text{Cov}(u_i, X_i) &= E(u_i X_i) - E(u_i)E(X_i) \\ &= E(X_i E(u_i|X_i)) - 0 \cdot E(X_i) \\ &= 0 \end{aligned}$$

where the law of iterated expectation is used twice at the second equality.

It follows that $\text{Cov}(u_i, X_i) \neq 0 \Rightarrow E(u_i|X_i) \neq 0$.

6.2 Assumption 2: (X_i, Y_i) for $i = 1, \dots, n$ are i.i.d.

- Each pair of X and Y , i.e., (X_i, Y_i) for $i = 1, \dots, n$, is selected randomly from the same joint distribution of X and Y .
- Since $u_i = Y_i - \beta_0 - \beta_1 X_i$, u_i is i.i.d., too.
- The cases that may violate of the i.i.d. assumption:
 - Time series data, $\text{Cov}(Y_t, Y_{t-1}) \neq 0$. That is, when we try to regress Y_t on Y_{t-1} , and if the current value Y_t depends on Y_{t-1} , which is very likely, the independence is violated. We call this violation as serial correlation.
 - Spatial data, $\text{Cov}(Y_r, Y_s) \neq 0$, where s and r refer to two neighboring regions. That is, when we try to regress Y_r on Y_s , they may well be correlated because they are adjacent. We call this violation as spatial correlation.

6.3 Assumption 3: large outliers are unlikely

$$0 < E(X_i^4) < \infty \text{ and } 0 < E(Y_i^4) < \infty$$

- A large outlier is an extreme value of X or Y .
- On a technical level, if X and Y are bounded, then they have finite fourth moments, i.e., finite kurtosis.
- The essence of this assumption is to say that a large outlier can strongly influence the results. So we need to rule out large outliers in estimation.

The influential observations and the leverage effects

- An outlier can be detected by a scatterplot. See Figure 5.
- There are also formal tests for the existence of the influential observations, some of which are coded in econometric software, like R and Stata.

7 Sampling Distribution of the OLS Estimators

7.1 Unbiasedness and consistency

The unbiasedness of $\hat{\beta}_0$ and $\hat{\beta}_1$

- The randomness of $\hat{\beta}_0$ and $\hat{\beta}_1$




Figure 5: How an outlier can influence the OLS estimates

Since (X_i, Y_i) for $i = 1, \dots, n$ are randomly drawn from a population, different draws can render different estimates, giving rise to the randomness of $\hat{\beta}_0$ and $\hat{\beta}_1$.

- The unbiasedness of $\hat{\beta}_0$ and $\hat{\beta}_1$

Let the true values of the intercept and the slope be β_0 and β_1 . Based on the least squares assumption #1: $E(u_i|X_i) = 0$

$$E(\hat{\beta}_0) = \beta_0 \text{ and } E(\hat{\beta}_1) = \beta_1$$

- Show that $\hat{\beta}_1$ is unbiased

Let's rewrite the formula of $\hat{\beta}_1$ here

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} \quad (20)$$

Given the random samples (X_i, Y_i) for $i = 1, \dots, n$, from $Y_i = \beta_0 + \beta_1 X_i + u_i$, We know that $\bar{Y} = \beta_0 + \beta_1 \bar{X} + \bar{u}$. It follows that $Y_i - \bar{Y} = \beta_1(X_i - \bar{X}) + u_i - \bar{u}$, Plugging it in the numerator in Equation (20). Then,

$$\begin{aligned} \sum_i (X_i - \bar{X})(Y_i - \bar{Y}) &= \sum_i (X_i - \bar{X}) [\beta_1(X_i - \bar{X}) + (u_i - \bar{u})] \\ &= \beta_1 \sum_i (X_i - \bar{X})^2 + \sum_i (X_i - \bar{X})u_i - \bar{u} \sum_i (X_i - \bar{X}) \\ &= \beta_1 \sum_i (X_i - \bar{X})^2 + \sum_i (X_i - \bar{X})u_i \end{aligned}$$

In the second equality, we use the fact that $\sum_i (X_i - \bar{X}) = 0$. Note that although we know from the first OLS assumption, $E(u_i) = 0$, we cannot guarantee that $\bar{u} = 0$ since u_1, \dots, u_n are simply random draws of u_i .

Substituting this expression in Equation (20) yields

$$\hat{\beta}_1 = \beta_1 + \frac{\frac{1}{n} \sum_i (X_i - \bar{X}) u_i}{\frac{1}{n} \sum_i (X_i - \bar{X})^2} \quad (21)$$

We prove that $\hat{\beta}_1$ is conditionally unbiased, from which the unconditional unbiasedness follows naturally.

$$\begin{aligned} E(\hat{\beta}_1 | X_1, \dots, X_n) &= \beta_1 + E \left\{ \left[\frac{\frac{1}{n} \sum_i (X_i - \bar{X}) u_i}{\frac{1}{n} \sum_i (X_i - \bar{X})^2} \right] \mid X_1, \dots, X_n \right\} \\ &= \beta_1 + \frac{\frac{1}{n} \sum_i (X_i - \bar{X}) E(u_i | X_1, \dots, X_n)}{\frac{1}{n} \sum_i (X_i - \bar{X})^2} \\ &= \beta_1 \quad (\text{by assumption 1}) \end{aligned}$$

It follows that

$$E(\hat{\beta}_1) = E(E(\hat{\beta}_1 | X_1, \dots, X_n)) = \beta_1$$

Therefore, $\hat{\beta}_1$ is an unbiased estimator of β_1 .

The proof of unbiasedness of $\hat{\beta}_0$ is left for exercise.

The consistency of $\hat{\beta}_0$ and $\hat{\beta}_1$

$\hat{\beta}$ is said to be a consistent estimator of β if as n goes to infinity, $\hat{\beta}$ is in probability close to β , which can be denoted as $n \rightarrow \infty, \hat{\beta} \xrightarrow{p} \beta$, or simply as $\text{plim}_{n \rightarrow \infty} \hat{\beta} = \beta$.

And the law of large number states that for random i.i.d. samples x_1, \dots, x_n , if $E(x_i) = \mu$ and $\text{Var}(x_i) < \infty$, then $\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i x_i = \mu$.

Then we can show that $\text{plim}_{n \rightarrow \infty} \hat{\beta}_1 = \beta_1$.

- A proof of consistency

The proof is not required to understand for this course. Therefore, you can skip it when you first read the notes.

From Equation (21) we can have

$$\text{plim}_{n \rightarrow \infty} (\hat{\beta}_1 - \beta_1) = \text{plim}_{n \rightarrow \infty} \frac{\frac{1}{n} \sum_i (X_i - \bar{X}) u_i}{\frac{1}{n} \sum_i (X_i - \bar{X})^2} = \frac{\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i (X_i - \bar{X}) u_i}{\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i (X_i - \bar{X})^2}$$

The denominator of the last equality is just a consistent estimator of the sample variance of X_i , that is, $\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i (X_i - \bar{X})^2 = \sigma_X^2$

Now we need to focus on $\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i (X_i - \bar{X})u_i$. To apply the law of large numbers, we need to find the expectation of $(X_i - \bar{X})u_i$. Given that $E(X_i u_i) = E(E(X_i u_i | X_i)) = E(X_i E(u_i | X_i)) = 0$, we have

$$E((X_i - \bar{X})u_i) = E(X_i u_i) + \frac{1}{n} \sum_i E(X_i u_i) = 0 + 0 = 0$$

So the variance of $(X_i - \bar{X})u_i$ can be expressed as

$$\begin{aligned} \text{Var}((X_i - \bar{X})u_i) &= E((X - \bar{X})^2 u_i^2) \\ &= E(E((X - \bar{X})^2 u_i^2 | X)) \\ &= E((X - \bar{X})^2 E(u_i^2 | X)) \\ &= E((X - \bar{X})^2 \sigma_u^2) \quad (\text{by the extended assumption 4. See Chapter 17}) \\ &< \infty \quad (\text{by assumption 3}) \end{aligned}$$

Since $E((X_i - \bar{X})u_i) = 0$, $\text{Var}((X_i - \bar{X})u_i) < \infty$, and X_i, u_i for $i = 1, \dots, n$ are i.i.d, by the law of large numbers, we have

$$\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i (X_i - \bar{X})u_i = 0$$

Therefore, $\text{plim}_{n \rightarrow \infty} \hat{\beta}_1 = \beta_1$.

Similarly, we can also prove that $\hat{\beta}_0$ is consistent, that is $\text{plim}_{n \rightarrow \infty} \hat{\beta}_0 = \beta_0$.

7.2 The asymptotic normal distribution

The central limit theory states that if X_1, \dots, X_n with the mean μ and the variance $0 < \sigma^2 < \infty$. Then, $\frac{1}{n} \sum_i X_i \xrightarrow{d} N(\mu, \frac{\sigma^2}{n})$.

From the proof of consistency, we have already seen that $E((X_i - \bar{X})u_i) = 0$, $\text{Var}((X_i - \bar{X})u_i) < \infty$, and X_i, u_i for $i = 1, \dots, n$ are i.i.d. By the central limit theory, we know that

$$\frac{1}{n} \sum_i (X_i - \bar{X})u_i \xrightarrow{d} N\left(0, \frac{1}{n} \text{Var}((X_i - \bar{X})u_i)\right)$$

It follows that from Equation (21) and the fact that $\text{plim}_{n \rightarrow \infty} \frac{1}{n} \sum_i (X_i - \bar{X})^2 = \text{Var}(X_i)$, $\hat{\beta}_1$ is asymptotically normally distributed as

$$\hat{\beta}_1 \xrightarrow{d} N\left(\beta_1, \sigma_{\hat{\beta}_1}^2\right)$$

where

$$\sigma_{\hat{\beta}_1}^2 = \frac{1}{n} \frac{\text{Var}((X_i - \bar{X})u_i)}{\text{Var}(X_i)^2} \quad (22)$$

Similarly, we can show that $\hat{\beta}_0 \xrightarrow{d} N(\beta_0, \sigma_{\hat{\beta}_0}^2)$, where

$$\sigma_{\hat{\beta}_0}^2 = \frac{1}{n} \frac{\text{Var}(H_i u_i)}{(E(H_i^2))^2}, \text{ and } H_i = 1 - \left(\frac{\mu_X}{E(X_i^2)} \right) X_i \quad (23)$$

- As $\text{Var}(X_i)$ increases, $\text{Var}(\hat{\beta}_1)$ decreases.
- As $\text{Var}(u_i)$ increases, $\text{Var}(\hat{\beta}_1)$ increases.




Figure 6: The Variance of $\hat{\beta}_1$ and the variance of X_i